

Assignment 2

Group Members

Abdul Wahab Madni, Hamza Badar, Aleksandr Semenikhin

```
In [44]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

sns.set_theme()
```

Ex.1: Kernel Density Estimation

```
In [2]: tdata = pd.read_table("values.txt", header=None, dtype=float, names=["Val",
tdata
```

Out[2]:

	Values
0	1.824378
1	3.446273
2	-0.427239
3	1.952763
4	1.107812
5	0.880768
6	3.022545
7	0.776415
8	-0.249058
9	1.802244
10	3.100371
11	-0.083352
12	1.969301
13	0.328054
14	0.996174
15	1.715757
16	3.271927
17	1.968713
18	1.458442

19 1.762478

20 3.000002

21 -0.168349

22 2.946189

23 1.235000

24 2.831157

25 -0.358687

26 -0.011253

27 0.850426

28 1.932236

29 1.001805

30 3.108471

31 2.051314

32 0.100576

33 -0.181802

34 3.074089

35 1.110291

36 1.008308


37 -0.211590

38 1.872469

39 2.513046

```
In [3]: sns.kdeplot(data= tdata)
```



```
Out[3]: <Axes: ylabel='Density'>
```


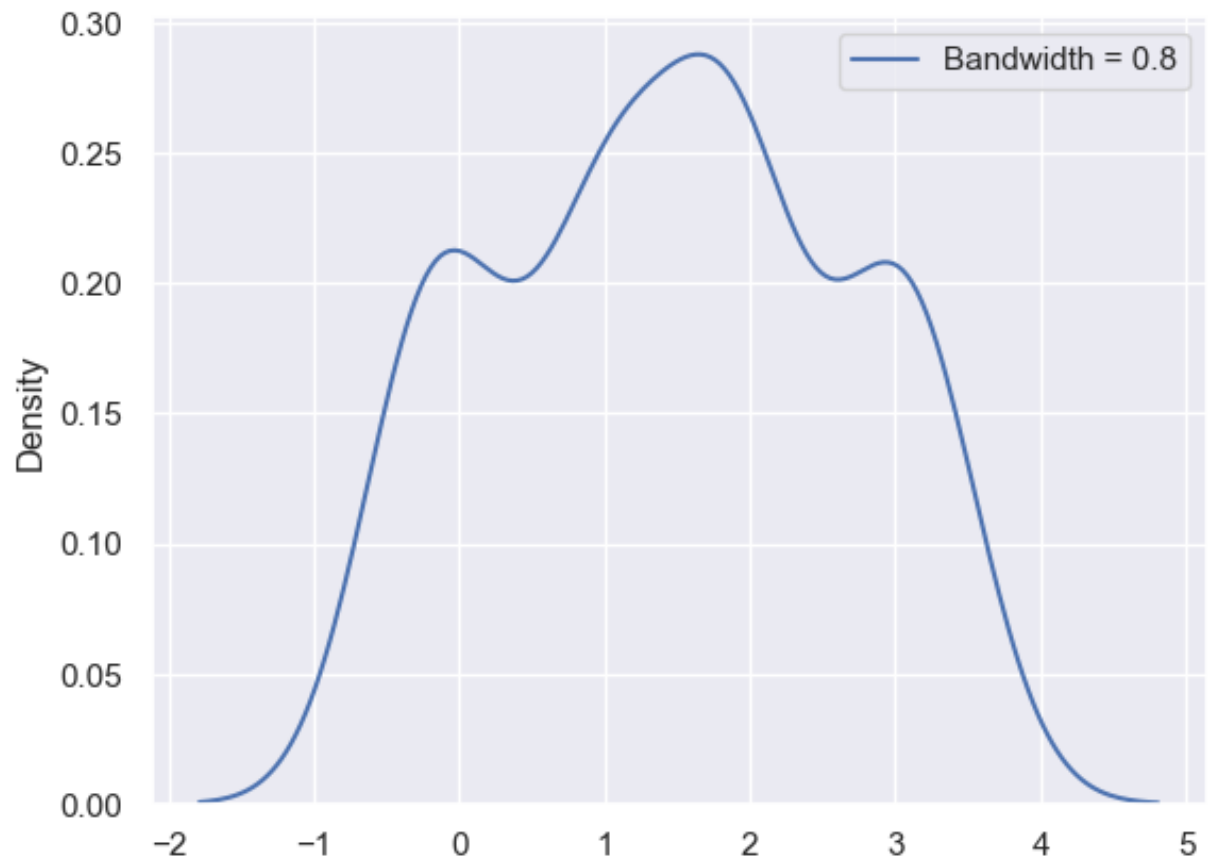



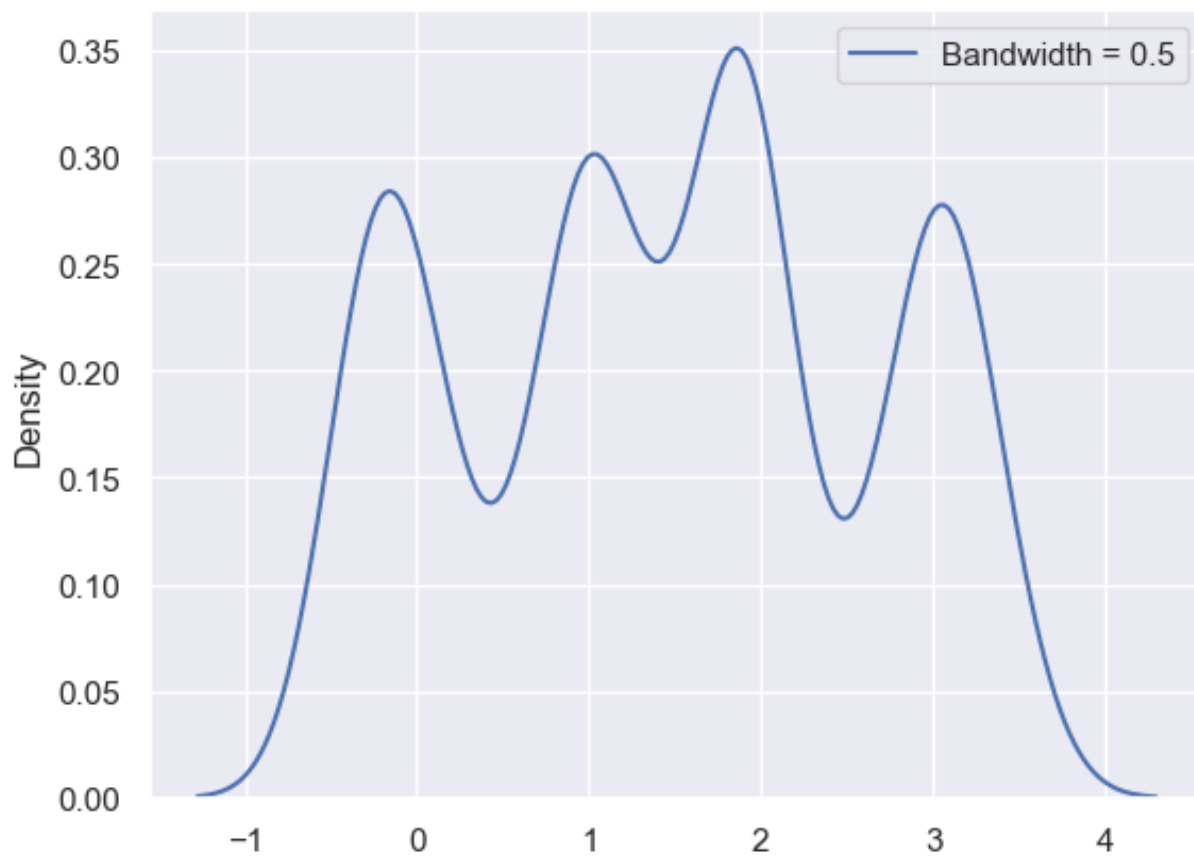
By visualization, we can see that the underlying distribution has a mode at around 1.5.


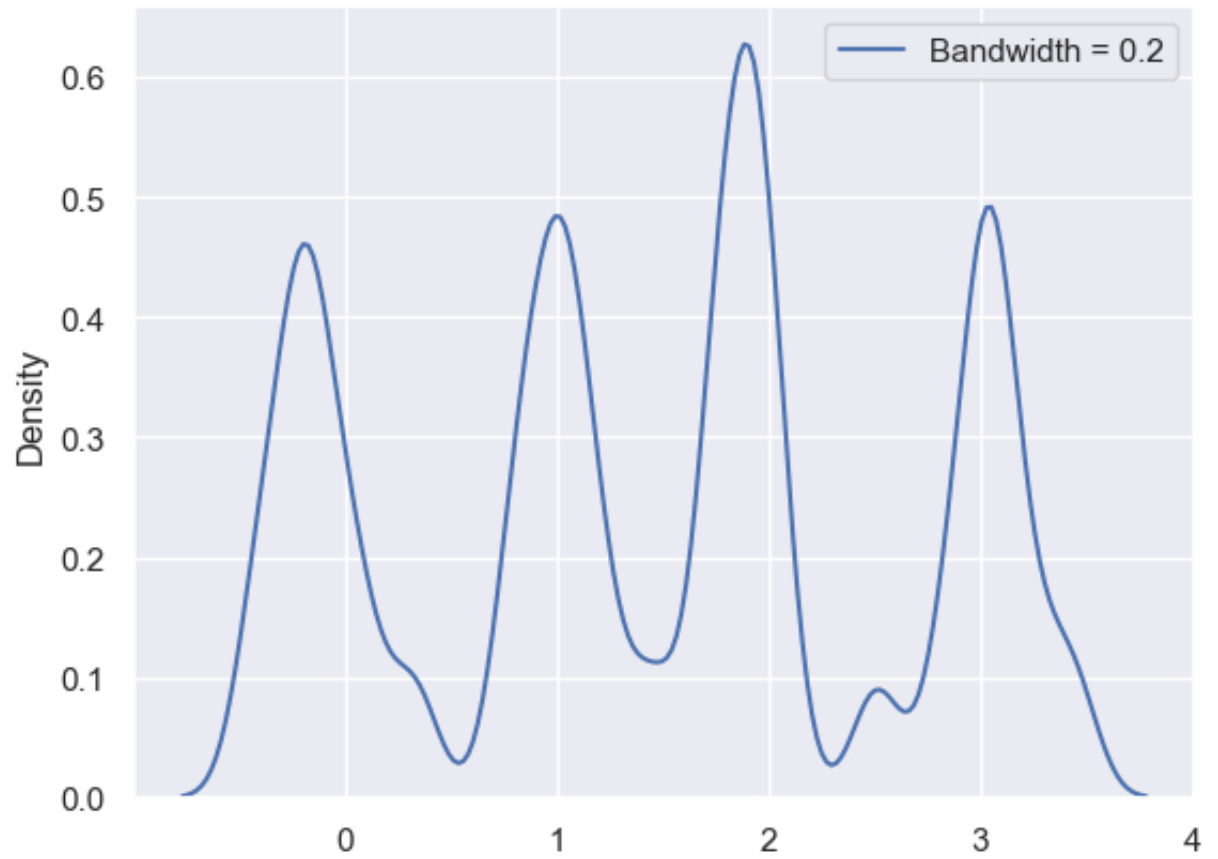
```
In [4]: bandwidths = [1.2, 1.0, 0.8, 0.6, 0.5, 0.4, 0.2, 0.1]

for i in bandwidths:
    plt.figure()
    sns.kdeplot(tdata, bw_adjust=i, label=f'Bandwidth = {i:.1f}')
    plt.legend()
```









Observations

From above KDEplots we can see that the bandwidth changes the shape of the estimated distribution greatly. As bandwidth value decreases, the distribution shape changes from the only one peak to four clear peaks (modes). The setting with 0.2 bandwidth shows the more precise distribution. From which we can observe that around -0.2, 1, 2, 3 it has the modes.

Ex. 2: Plotting Categorical Data

```
In [5]: df2 = pd.read_excel('chronic_kidney_disease_numerical.xls')
```

```
In [6]: df2
```

```
Out[6]:
```

	age	blood pressure	specific gravity	albumin	sugar	blood glucose random	blood urea	serum creatinine	sodium	potas
0	48.0	80.0	1.020	1.0	0.0	121.0	36.0	1.2	NaN	
1	7.0	50.0	1.020	4.0	0.0	NaN	18.0	0.8	NaN	
2	62.0	80.0	1.010	2.0	3.0	423.0	53.0	1.8	NaN	
3	48.0	70.0	1.005	4.0	0.0	117.0	56.0	3.8	111.0	
4	51.0	80.0	1.010	2.0	0.0	106.0	26.0	1.4	NaN	
...
395	55.0	80.0	1.020	0.0	0.0	140.0	49.0	0.5	150.0	
396	42.0	70.0	1.025	0.0	0.0	75.0	31.0	1.2	141.0	
397	12.0	80.0	1.020	0.0	0.0	100.0	26.0	0.6	137.0	
398	17.0	60.0	1.025	0.0	0.0	114.0	50.0	1.0	135.0	
399	58.0	80.0	1.025	0.0	0.0	131.0	18.0	1.1	141.0	

400 rows × 15 columns

```
In [8]: lm = pd.melt(df2, id_vars=['class'])
lm
```


Out[8]:

	class	variable	value
0	ckd	age	48.0
1	ckd	age	7.0
2	ckd	age	62.0
3	ckd	age	48.0
4	ckd	age	51.0
...
5595	notckd	red blood cell count	4.9
5596	notckd	red blood cell count	6.2
5597	notckd	red blood cell count	5.4
5598	notckd	red blood cell count	5.9
5599	notckd	red blood cell count	6.1


5600 rows × 3 columns

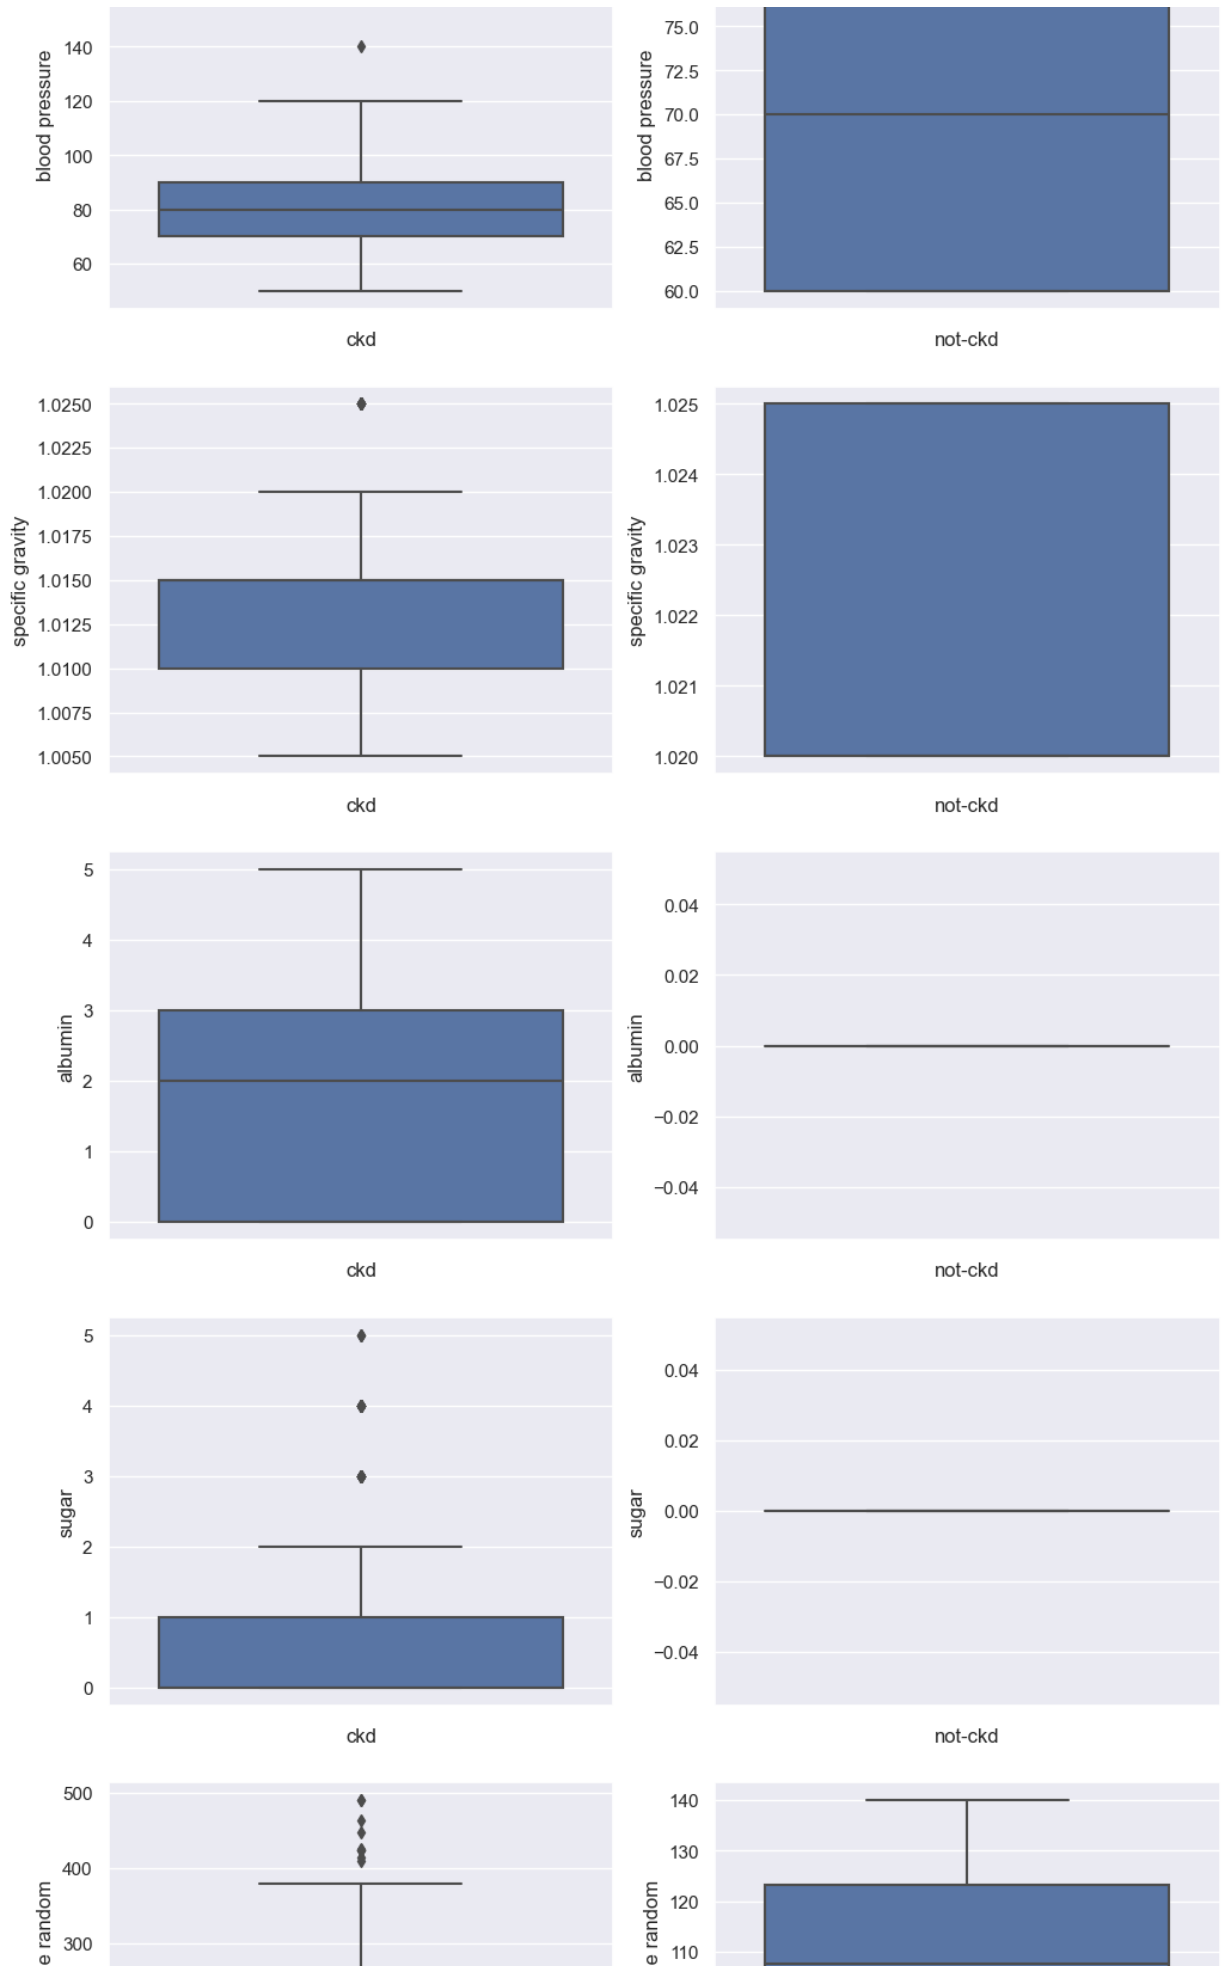
To disable sharing y axes between the two attributes, we made two groups of data on the basis of chronic kidney disease.

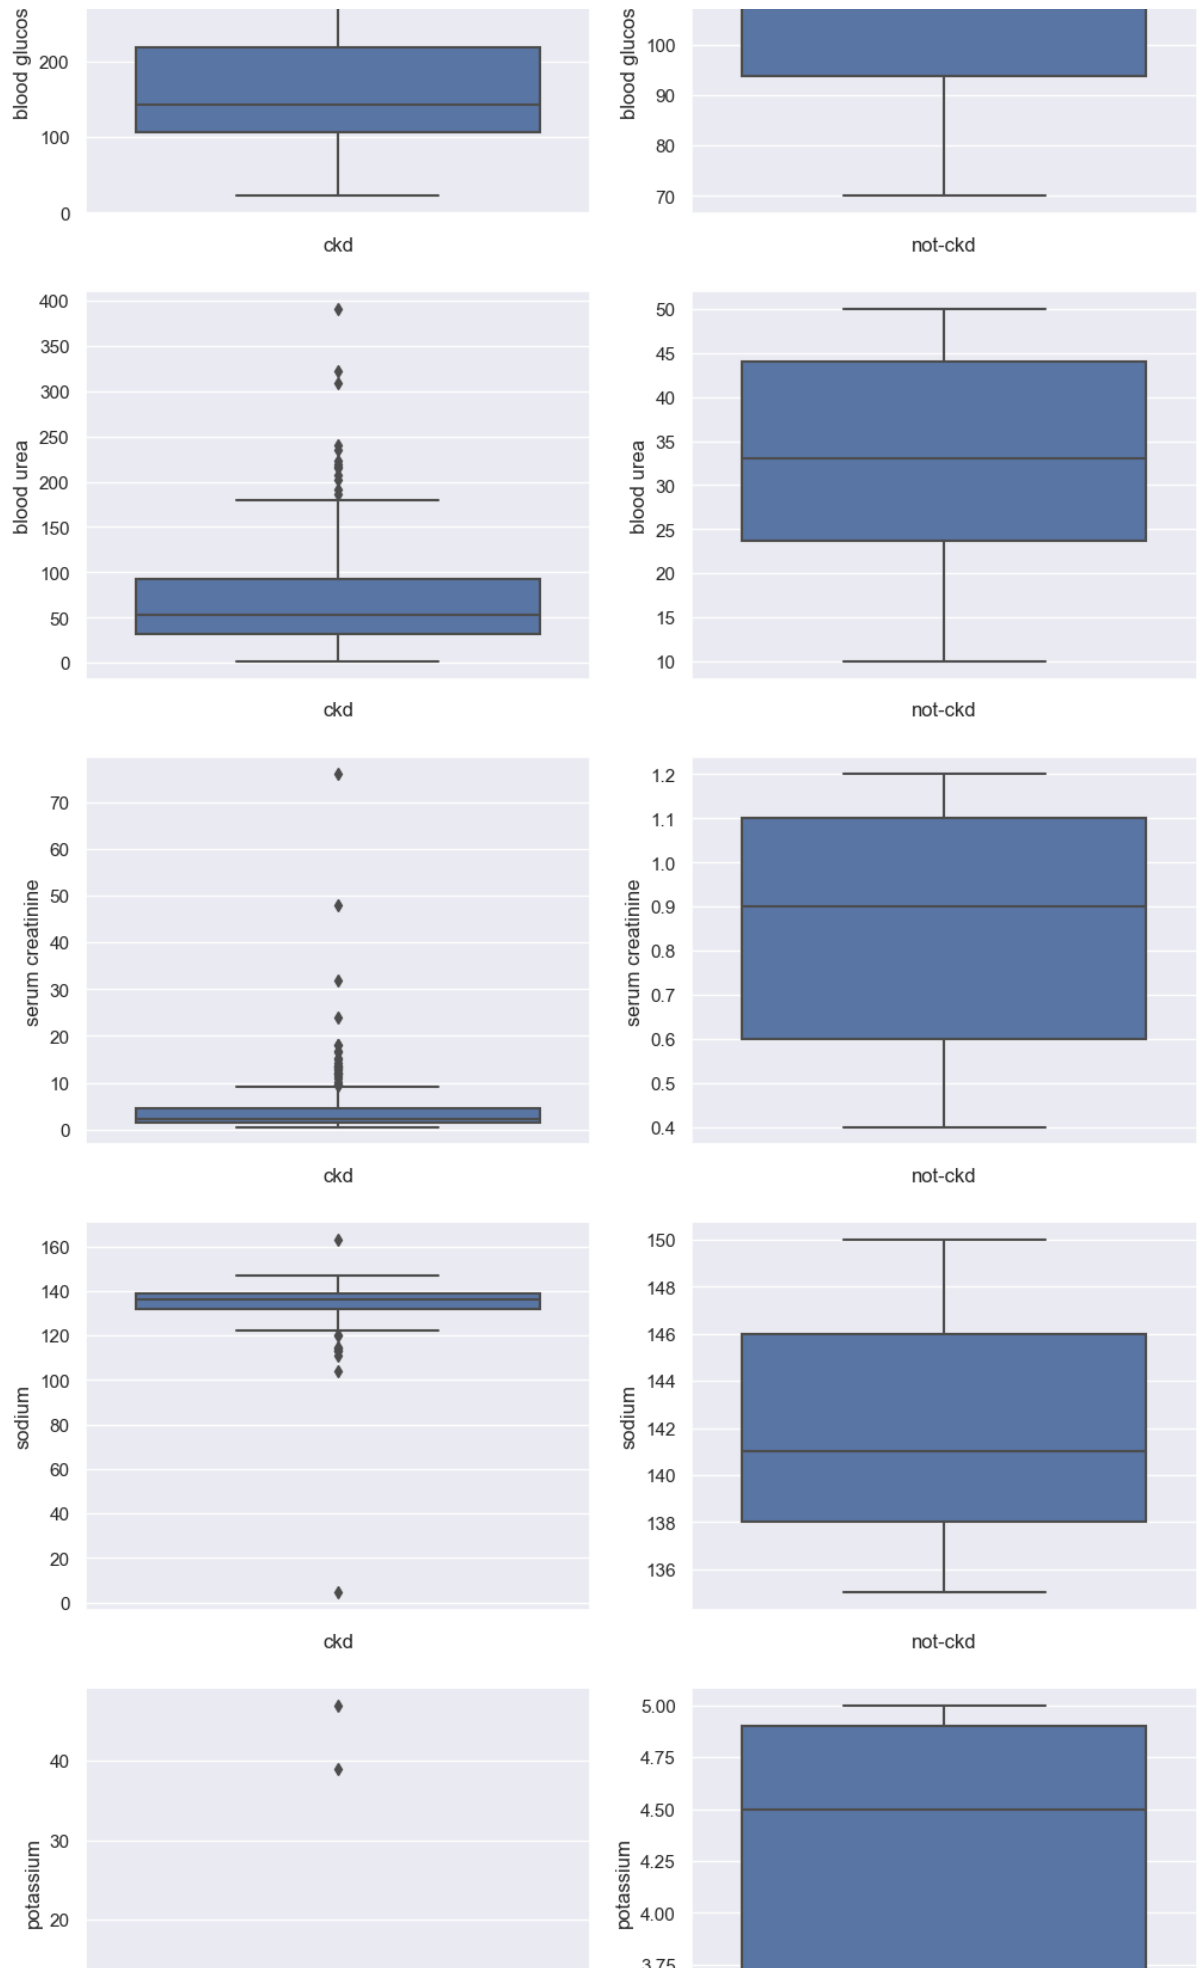
```
In [9]: ckd_group = df2[df2['class'] == 'ckd']
non_ckd_group = df2[df2['class'] == 'notckd']
#for passing only the float and int datatype
num_attributes = [col for col in df2.columns if df2[col].dtype == 'float64']

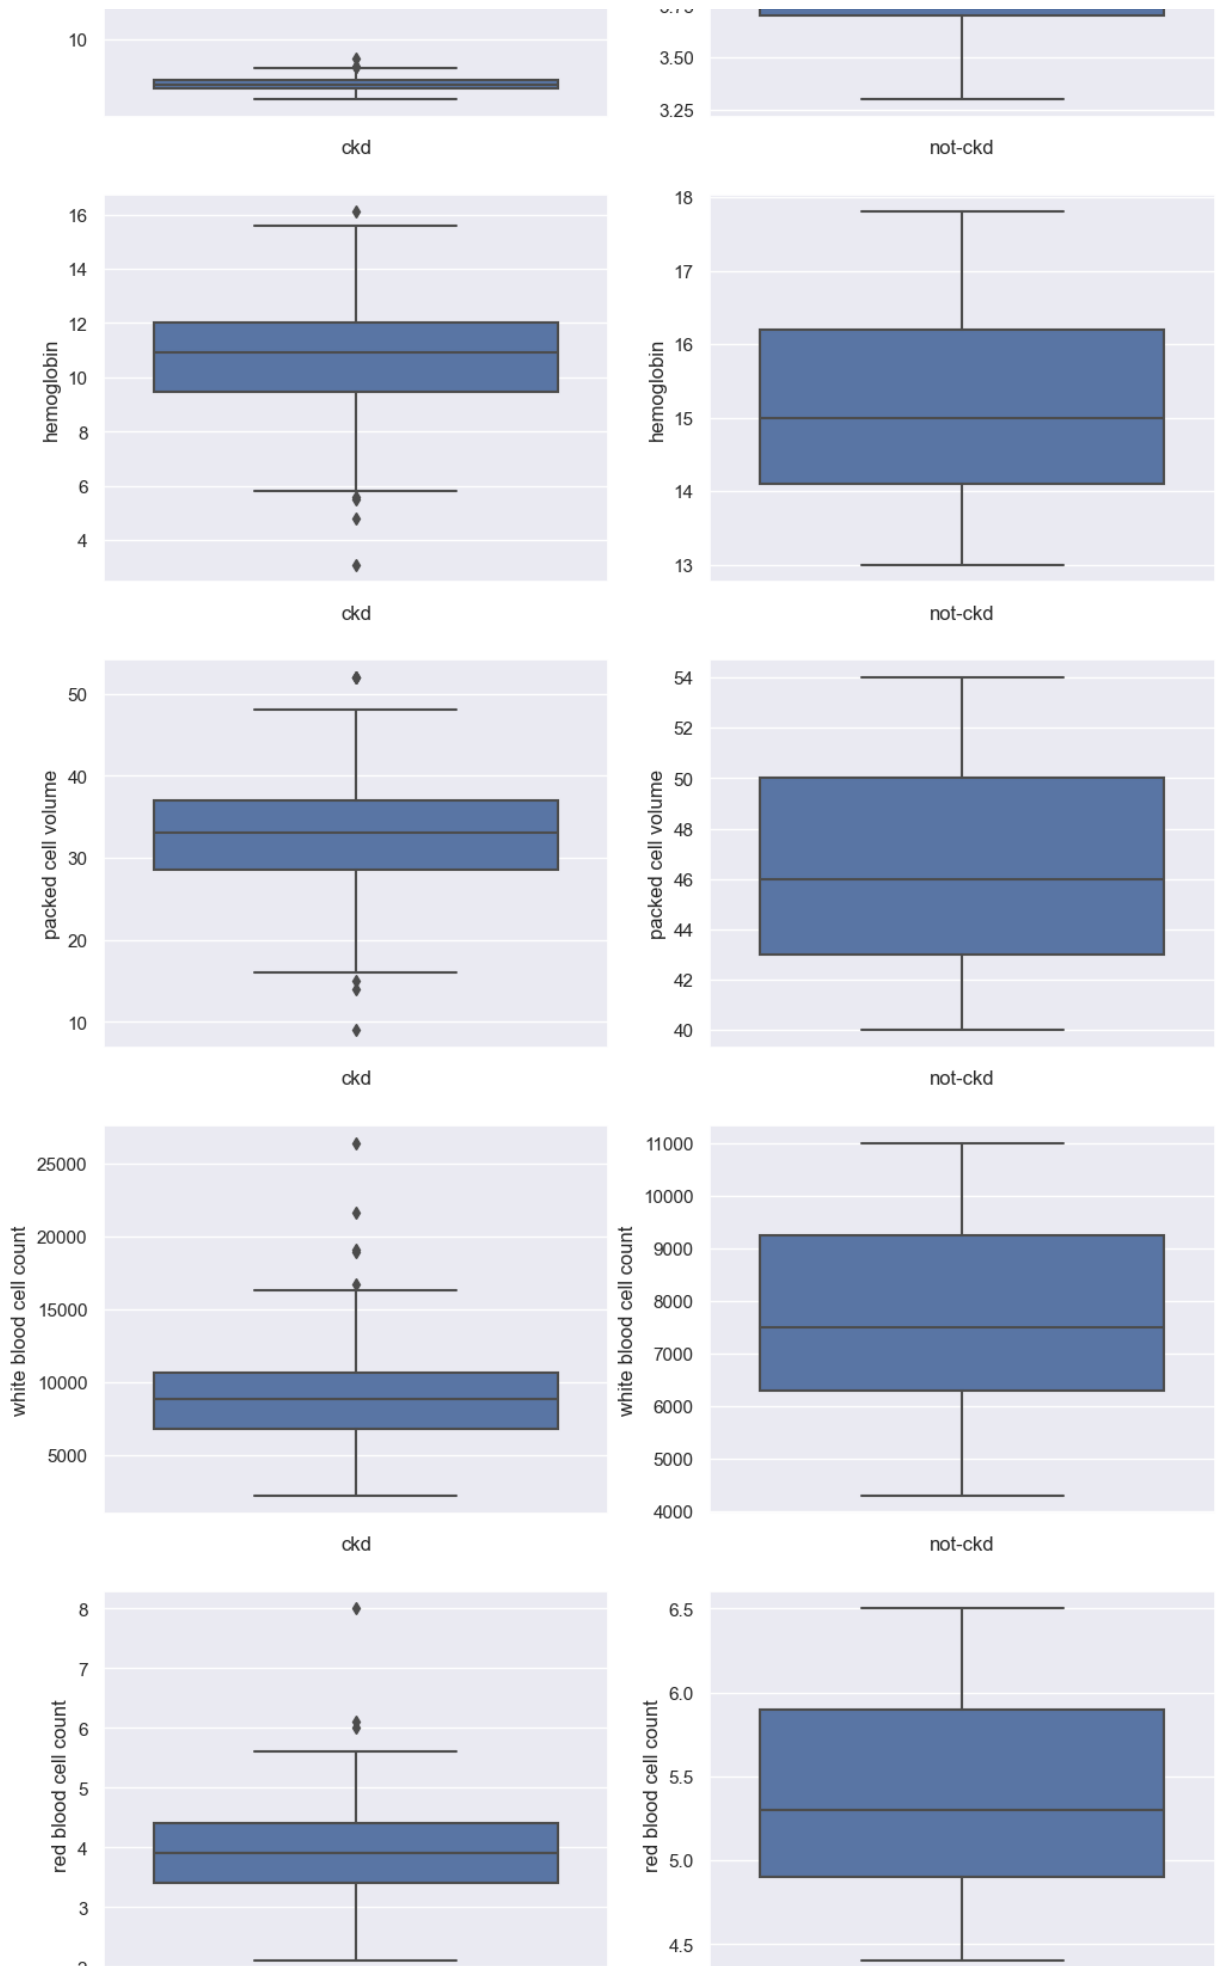
fig, axes = plt.subplots(nrows=len(num_attributes), ncols=2, figsize=(12,

for i, value in enumerate(num_attributes):
    sns.boxplot(y=value, data=ckd_group, ax=axes[i, 0])
    sns.boxplot(y=value, data=non_ckd_group, ax=axes[i, 1])
    axes[i, 0].set_xlabel('ckd')
    axes[i, 1].set_xlabel('not-ckd')
```









ckd

not-ckd

By visualization, we can see that the 'albumin' attribute is a high indicator of chronic kidney disease. Whereas the 'potassium' seems to be unrelated to it.

Ex.3: Multidimensional data filtering and visualization


```
In [10]: df = pd.read_csv("winequality-red.csv", sep=';')
df.head(10)
```

```
Out[10]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56
5	7.4	0.66	0.00	1.8	0.075	13.0	40.0	0.9978	3.51	0.56
6	7.9	0.60	0.06	1.6	0.069	15.0	59.0	0.9964	3.30	0.46
7	7.3	0.65	0.00	1.2	0.065	15.0	21.0	0.9946	3.39	0.47
8	7.8	0.58	0.02	2.0	0.073	9.0	18.0	0.9968	3.36	0.57
9	7.5	0.50	0.36	6.1	0.071	17.0	102.0	0.9978	3.35	0.80

```
In [11]: sns.histplot(df, x="quality")
```

```
Out[11]: <Axes: xlabel='quality', ylabel='Count'>
```



```
In [12]: wine_quality = {3:"low",4:"low",5:"medium",6:"medium",7:"high",8:"high"}  
df_renamed = df.rename(columns={'quality':'quality bin'})
```

```
In [13]: df_renamed['quality bin'] = df_renamed['quality bin'].map(wine_quality)
```

```
In [14]: df_renamed["quality bin"].head(20)
```

```
Out[14]: 0    medium  
1    medium  
2    medium  
3    medium  
4    medium  
5    medium  
6    medium  
7     high  
8     high  
9    medium  
10   medium  
11   medium  
12   medium  
13   medium  
14   medium  
15   medium  
16   high  
17   medium  
18    low  
19   medium  
Name: quality bin, dtype: object
```

```
In [15]: df_renamed.head(10)
```

```
Out[15]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56
5	7.4	0.66	0.00	1.8	0.075	13.0	40.0	0.9978	3.51	0.56
6	7.9	0.60	0.06	1.6	0.069	15.0	59.0	0.9964	3.30	0.46
7	7.3	0.65	0.00	1.2	0.065	15.0	21.0	0.9946	3.39	0.47
8	7.8	0.58	0.02	2.0	0.073	9.0	18.0	0.9968	3.36	0.57
9	7.5	0.50	0.36	6.1	0.071	17.0	102.0	0.9978	3.35	0.80


```
In [16]: df_quality_filter = df_renamed[(df_renamed["quality bin"] == "low") | (df
```

```
In [24]: df_quality_filter.head(20)
```

Out[24]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphat
7	7.3	0.650	0.00	1.2	0.065	15.0	21.0	0.99460	3.39	0.
8	7.8	0.580	0.02	2.0	0.073	9.0	18.0	0.99680	3.36	0.
16	8.5	0.280	0.56	1.8	0.092	35.0	103.0	0.99690	3.30	0.
18	7.4	0.590	0.08	4.4	0.086	6.0	29.0	0.99740	3.38	0.
37	8.1	0.380	0.28	2.1	0.066	13.0	30.0	0.99680	3.23	0.
38	5.7	1.130	0.09	1.5	0.172	7.0	19.0	0.99400	3.50	0.
41	8.8	0.610	0.30	2.8	0.088	17.0	46.0	0.99760	3.26	0
45	4.6	0.520	0.15	2.1	0.054	8.0	65.0	0.99340	3.90	0.
62	7.5	0.520	0.16	1.9	0.085	12.0	35.0	0.99680	3.38	0.
73	8.3	0.675	0.26	2.1	0.084	11.0	43.0	0.99760	3.31	0.
79	8.3	0.625	0.20	1.5	0.080	27.0	119.0	0.99720	3.16	1
94	5.0	1.020	0.04	1.4	0.045	41.0	85.0	0.99380	3.75	0.
128	8.0	0.590	0.16	1.8	0.065	3.0	16.0	0.99620	3.42	0.
151	9.2	0.520	1.00	3.4	0.610	32.0	69.0	0.99960	2.74	2.
161	7.6	0.680	0.02	1.3	0.072	9.0	20.0	0.99650	3.17	1.
167	7.3	0.550	0.03	1.6	0.072	17.0	42.0	0.99560	3.37	0.
170	7.9	0.885	0.03	1.8	0.058	4.0	8.0	0.99720	3.36	0.
198	5.4	0.835	0.08	1.2	0.046	13.0	93.0	0.99240	3.57	0.
199	6.9	1.090	0.06	2.1	0.061	12.0	31.0	0.99480	3.51	0.
200	9.6	0.320	0.47	1.4	0.056	9.0	24.0	0.99695	3.22	0.

In [25]: `sns.pairplot(data=df_quality_filter, hue="quality bin",diag_kind="hist")`Out[25]: `<seaborn.axisgrid.PairGrid at 0x17f7bb8b0>`



list of attributes that distinguish high quality out of low quality wine:

1. Volatile Acidity
2. Citric Acid
3. PH level
4. Sulphates
5. Alcohol

```
In [26]: from sklearn.feature_selection import SelectKBest
         from sklearn.feature_selection import f_classif
```

```
In [27]: X = df_quality_filter.iloc[:,0:11]
         y = df_quality_filter.iloc[:,11]
```

```
In [28]: selector = SelectKBest(f_classif, k=5)
         selector.fit(X,y)
         cols = selector.get_support(indices=True)
         df_focused = df_quality_filter.iloc[:,cols]
```

```
In [29]: df_focused.head(10)
```

```
Out[29]:
```

	volatile acidity	citric acid	pH	sulphates	alcohol
7	0.650	0.00	3.39	0.47	10.0
8	0.580	0.02	3.36	0.57	9.5
16	0.280	0.56	3.30	0.75	10.5
18	0.590	0.08	3.38	0.50	9.0
37	0.380	0.28	3.23	0.73	9.7
38	1.130	0.09	3.50	0.48	9.8
41	0.610	0.30	3.26	0.51	9.3
45	0.520	0.15	3.90	0.56	13.1
62	0.520	0.16	3.38	0.62	9.5
73	0.675	0.26	3.31	0.53	9.2

```
In [41]: df_focused = df_focused.copy()
df_focused["quality bin"] = df_quality_filter['quality bin']
```

```
In [42]: df_focused
```

```
Out[42]:
```

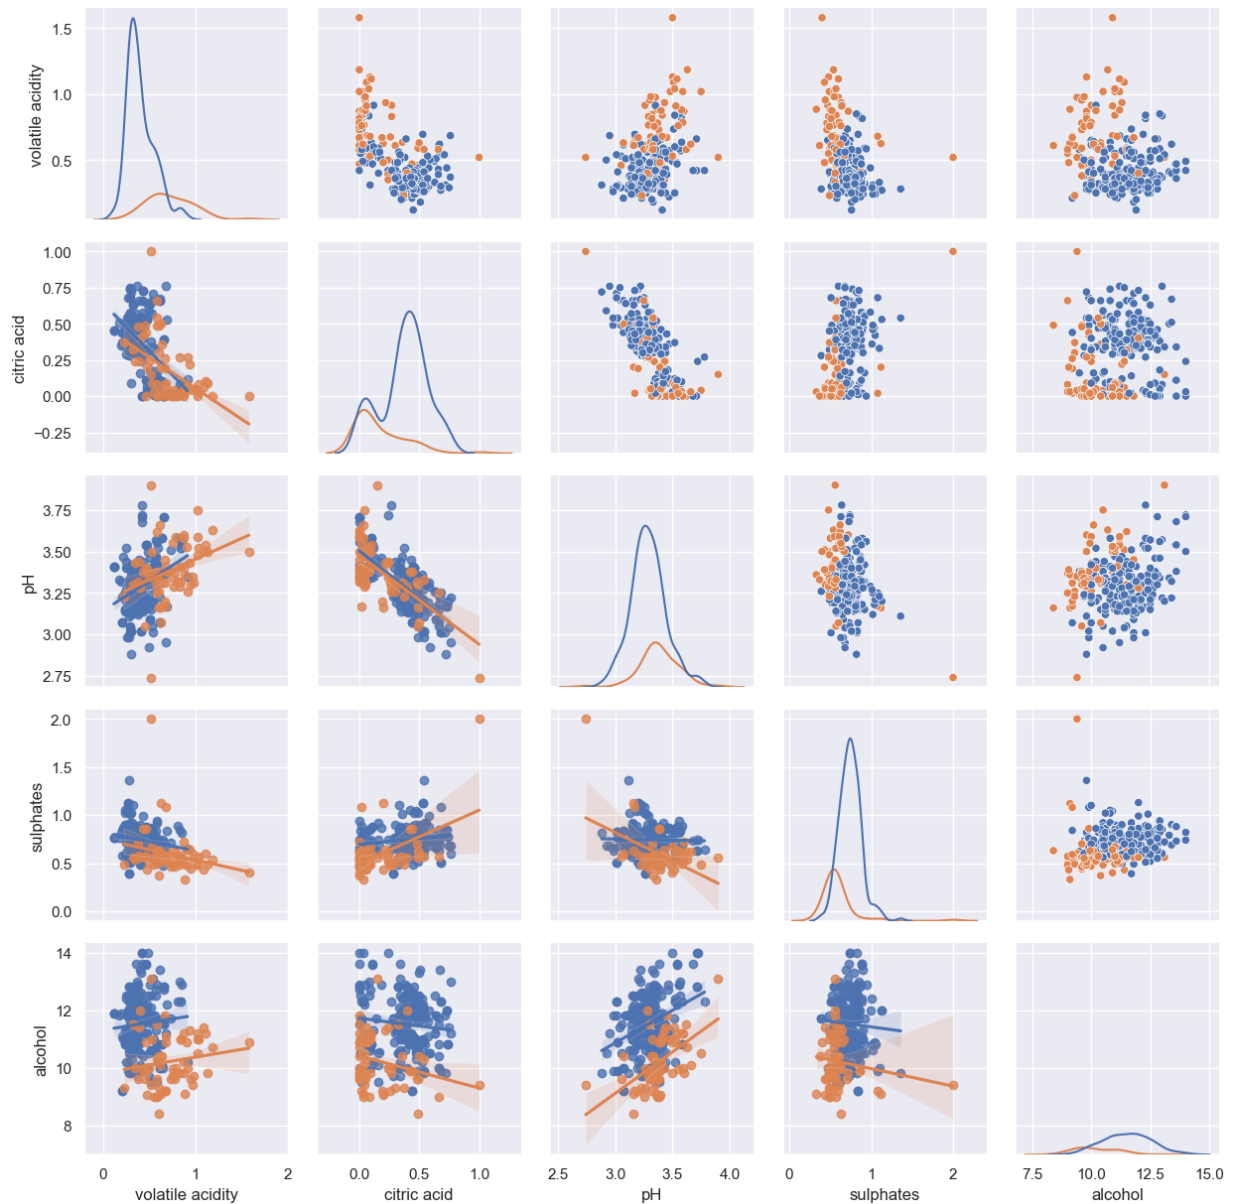
	volatile acidity	citric acid	pH	sulphates	alcohol	quality bin
7	0.65	0.00	3.39	0.47	10.00	high
8	0.58	0.02	3.36	0.57	9.50	high
16	0.28	0.56	3.30	0.75	10.50	high
18	0.59	0.08	3.38	0.50	9.00	low
37	0.38	0.28	3.23	0.73	9.70	high
...
1541	0.25	0.29	3.40	0.76	10.90	high
1544	0.37	0.43	3.17	0.81	11.20	high
1549	0.36	0.30	3.24	0.70	11.40	high
1555	0.56	0.17	3.44	0.68	10.55	high
1584	0.32	0.44	3.29	0.80	11.60	high

280 rows × 6 columns

In total, we can see, that attributes, that were chosen, are identical with those after feature selection

```
In [32]: g = sns.PairGrid(df_focused, hue="quality bin")
g.map_diag(sns.kdeplot)
g.map_upper(sns.scatterplot)
g.map_lower(sns.regplot)
```

```
Out[32]: <seaborn.axisgrid.PairGrid at 0x284205960>
```



Regardless of quality of wine we have Alcohol percentage, PH level and Citric Acid that have strong correlation. At the same time there is Sulphates that has different correlation depending on the wine quality. To my mind, citric acid definitely has multimodal distribution. Moreover, I think that Alcohol percentage can also have multimodal distribution. We can easily see four outliers all in low quality wine. It's:

1. 1.580 volatile acidity
2. 1.0 citric acid
3. 2.74 pH
4. 2.0 sulphates

```
In [33]: from numpy import sqrt

def distanceCalc(xA,xB,yA,yB):
    return sqrt(pow((xA-xB),2) + pow((yA-yB),2))
```

```
In [34]: def DSC(highDSC,lowDSC,clustersPoints):
    dscL = dscH = 0
    for i in range(len(highDSC)):
        if (distanceCalc(clustersPoints[0][0],highDSC[i][0],clustersPoint
            dscH += 1
            # print(dscH)
    # print("/")
    for i in range(len(lowDSC)):
        if (distanceCalc(clustersPoints[1][0],lowDSC[i][0],clustersPoints
            dscL += 1
            # print(dscL)
    return {(dscH/len(highDSC))*100,(dscL/len(lowDSC))*100}
```

```
In [35]: clusterCenters = df_focused.groupby(["quality bin"]).mean()[ :2].values
```

```
In [36]: labels = df_focused.columns.values
lowPoints = df_focused.loc[df_focused["quality bin"] == "low", labels[:5]]
highPoints = df_focused.loc[df_focused["quality bin"] == "high", labels[:5]]
```

```
In [37]: for i in range(1,5):
    lowPair = [[]]* len(lowPoints)
    highPair = [[]]* len(highPoints)
    for j in range(len(lowPoints)):
        lowPair[j] = [lowPoints[j][0],lowPoints[j][i]]
    for g in range(len(highPoints)):
        highPair[g] = [highPoints[g][0],highPoints[g][i]]
    firstClusters = [[clusterCenters[0][0],clusterCenters[0][i]],cluster
    print("pair of: Volatile acidity and ",i,DSC(highPair, lowPair, first
```

```
pair of: Volatile acidity and 1 {26.984126984126984, 20.737327188940093}
pair of: Volatile acidity and 2 {16.129032258064516, 26.984126984126984}
pair of: Volatile acidity and 3 {14.285714285714285, 22.22222222222222}
pair of: Volatile acidity and 4 {25.806451612903224, 30.158730158730158}
```

```
In [38]: for i in range(2,5):
    lowPair = [[]]* len(lowPoints)
    highPair = [[]]* len(highPoints)
    for j in range(len(lowPoints)):
        lowPair[j] = [lowPoints[j][1],lowPoints[j][i]]
    for g in range(len(highPoints)):
        highPair[g] = [highPoints[g][1],highPoints[g][i]]
    firstClusters = [[clusterCenters[0][1],clusterCenters[0][i]],cluster
    print("pair of: Citric Acid and ",i,DSC(highPair, lowPair, firstClust
```

```
pair of: Citric Acid and 2 {28.57142857142857, 22.58064516129032}
pair of: Citric Acid and 3 {22.22222222222222, 23.04147465437788}
pair of: Citric Acid and 4 {25.806451612903224, 28.57142857142857}
```

```
In [110]: for i in range(3,5):
            lowPair = [[]]* len(lowPoints)
            highPair = [[]]* len(highPoints)
            for j in range(len(lowPoints)):
                lowPair[j] = [lowPoints[j][2],lowPoints[j][i]]
            for g in range(len(highPoints)):
                highPair[g] = [highPoints[g][2],highPoints[g][i]]
            firstClusters = [[clusterCenters[0][2],clusterCenters[0][i]],clusterCenters[0][i]]
            print("pair of: pH and ",i,DSC(highPair, lowPair, firstClusters))

pair of: pH and 3 {21.658986175115206, 15.873015873015872}
pair of: pH and 4 {33.33333333333333, 26.72811059907834}
```

```
In [39]: for i in range(4,5):
            lowPair = [[]]* len(lowPoints)
            highPair = [[]]* len(highPoints)
            for j in range(len(lowPoints)):
                lowPair[j] = [lowPoints[j][3],lowPoints[j][i]]
            for g in range(len(highPoints)):
                highPair[g] = [highPoints[g][3],highPoints[g][i]]
            firstClusters = [[clusterCenters[0][3],clusterCenters[0][i]],clusterCenters[0][i]]
            print("pair of: Sulphates and ",i,DSC(highPair, lowPair, firstClusters))

pair of: Sulphates and 4 {33.33333333333333, 26.72811059907834}
```

So in total we can highlight 3 main pairs of arguments:

1. Volatile acidity and Alcohol
2. pH and Alcohol
3. Sulphates and Alcohol

Ex.4: Evaluating PCP Variants

a)

in addition to standard PCP, the authors tested the following eight variations:

1. In the first variation, they added scatter plots into standard PCP. They further rotated the scatter plots by 45° on axis labels. They expected that the embedding of the scatter plots to PCP will make it more intuitive to visualize.
2. In the second variant, they added density based coloring to the polylines of the standard PCP using lookup table(LUT). They expected it to improve visualisation because of the strong visual cue of colors.
3. They changed opacity with respect to density of the polylines to achieve blending. Authors expected it to reduce the noise as noise reduction will make the pcps easy to comprehend.
4. In this variation, authors combined previous two(color, blending) techniques together. According to them ColorBlending should have the properties of both variations hence will prove to be a more better variant.
5. This variation replaces straight lines with curves to improve the clarity and smoothness of the visualization. They expected it to remove ambiguities.
6. Random Tour variation is a rotation of PCPs by animation which shows different angles using time dependent matrix. Authors expected it to perform not as well as non-animated variations because even though it shows new patterns of data to the viewer, it is still difficult to comprehend.
7. Permutation tour is an alternative animation scheme which smoothly cycles through all the permutations, showing all possible axis in an animated fashion. For this, they had same expectations as Random Tour because both of these are animated techniques.
8. In Wobble Variation, polylines wobble around a fixed base position in a sinusoidal way and keeps clusters fairly fixed. Because of its fixed base position, authors expected it to perform better than previous animated techniques.

b)

Authors did not expected the Color, Blending and ColorBlending variations to perform that badly. To their surprise, scatter plots regardless of their simplicity outperformed every other variation. They were also wrong with placing the standard pcg at last position as it performed 2nd best. However, they guessed quite accurately about the animation related variations(random tour, permutation tour, wobble).

c)

If we had to design a visualization with 5 or less dimensions, we'll prefer the color blend variation as it makes the pcg very clear to visualize as long as you have less than 6 dimensions. otherwise, we'll prefer the scatterplot variation as it turned out to be more understandable to the viewers as the findings of Li et al. [LMvW08] and this paper clearly indicates.